Univariate Time Series Data and Model Card

Generated by Cardtale - Automated Model and Data Card Generator

This report provides an automated, comprehensive analysis of univariate time series data. Generated by Cardtale, it explores basic aspects and potential challenges in your data to support informed decision-making and modeling choices.

Generated: 2024-10-29 10:58

Series Name: Q1

Table of Contents

Other aspects were explored but omitted from the final report:

Data Overview

This section examines the core characteristics and statistical properties of the time series. Understanding these attributes is important for assessing data quality and gaining a preliminary context. We explore the temporal structure, summary statistics, and distribution patterns to create a baseline understanding of your data.

Time Series Plot

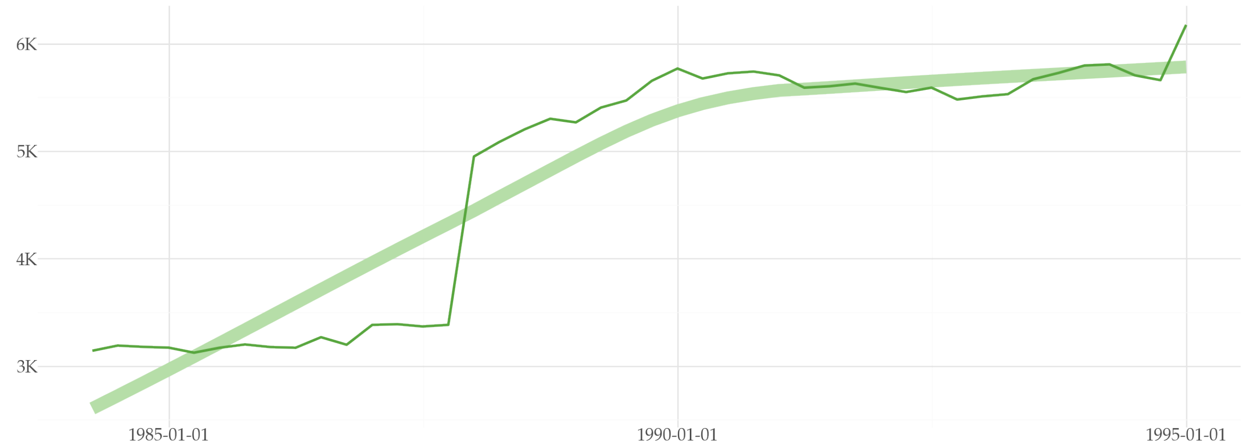

Figure 1: Time series line plot.

- A total of 44 quarterly observations which span from 1984-13 to 1994-52.
- The mean value of the series is 4775.59 (median equal to 5477.05), with a standard deviation of 1139.37. The data ranges from a minimum of 3124.38 to a maximum of 6176.6.

Data Distribution

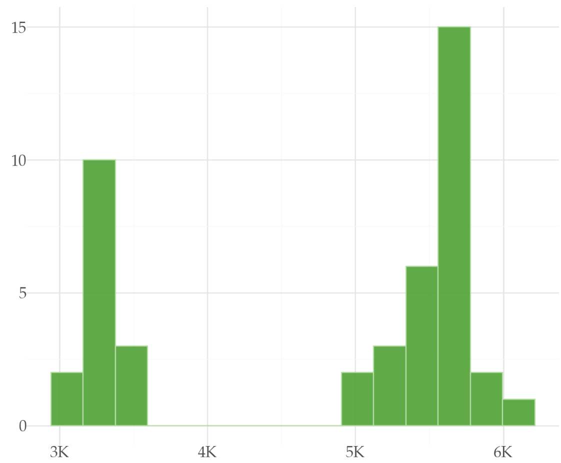

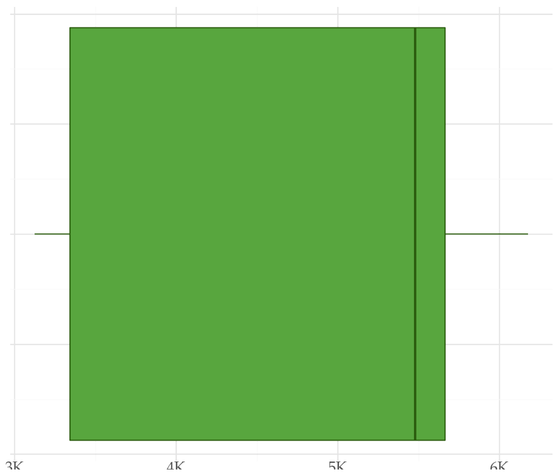

Figure 2: Distribution of the time series using an histogram (left) and a boxplot (right).

- The Kolmogorov-Smirnov test rejects the hypothesis that the series is distributed according to the following distributions: Logistic, Gaussian, Exponentially Modified Gaussian distribution, Chi-squared, Gamma, Cauchy, Exponential, Pareto, Log-Normal, and Power-law
- The statistical tests carried out did not find any suitable distribution for the data with an acceptable significance.

- There are no outliers in the data set according to the boxplot representation.
- The excess kurtosis is equal to -1.53. This indicates that the data has a light tailed distribution.
- The skewness is equal to -0.59, which is close to zero. This indicates a symmetric distribution, though there is a slight left skewness.

Trend and Seasonality

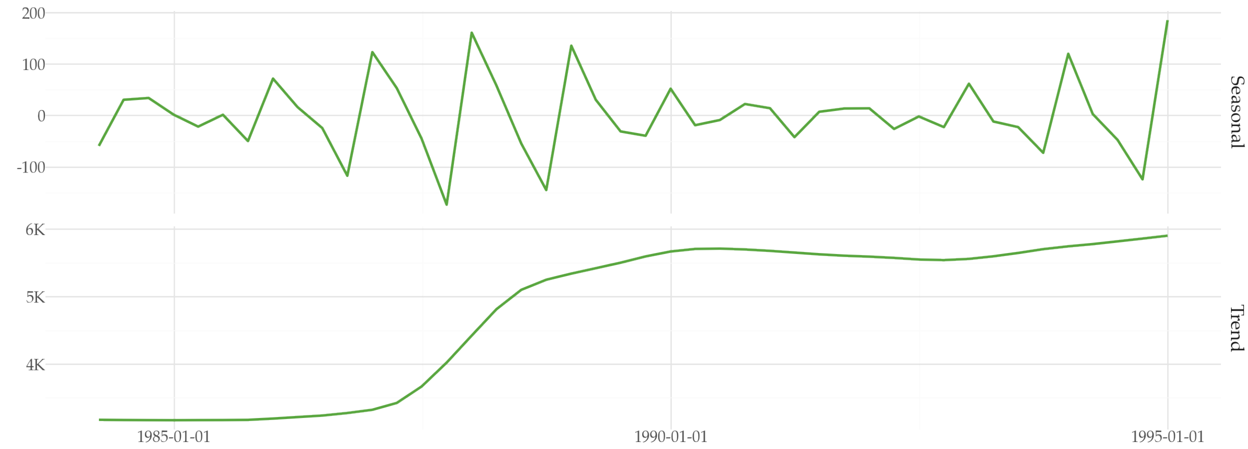

Figure 3: Seasonal and Trend components after decomposition using the STL (Season-Trend decomposition using LOESS) method.

- All hypothesis tests carried out (KPSS, Augmented Dickey-Fuller, and Philips-Perron) indicate that the time series is not stationary in trend/level.
- All hypothesis tests carried out (Wang-Smith-Hyndman and OCSB) indicate that the time series is not stationary in seasonality for the specified period.

Auto-Correlation

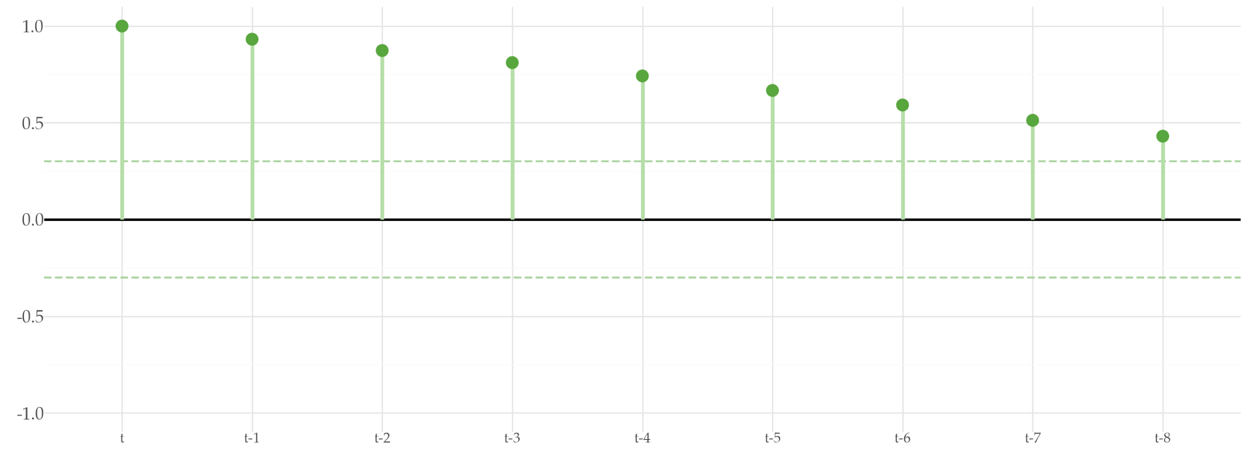

Figure 4: Auto-correlation plot up to 8 lags.

- The following lags show significant autocorrelation: t-1, t-2, t-3, t-4, t-5, t-6, t-7, and t-8. The autocorrelation is positive for all lags with a significant value.

- All lags relative to the seasonal period (t-4 and t-8) show a significant autocorrelation.

Partial Auto-Correlation

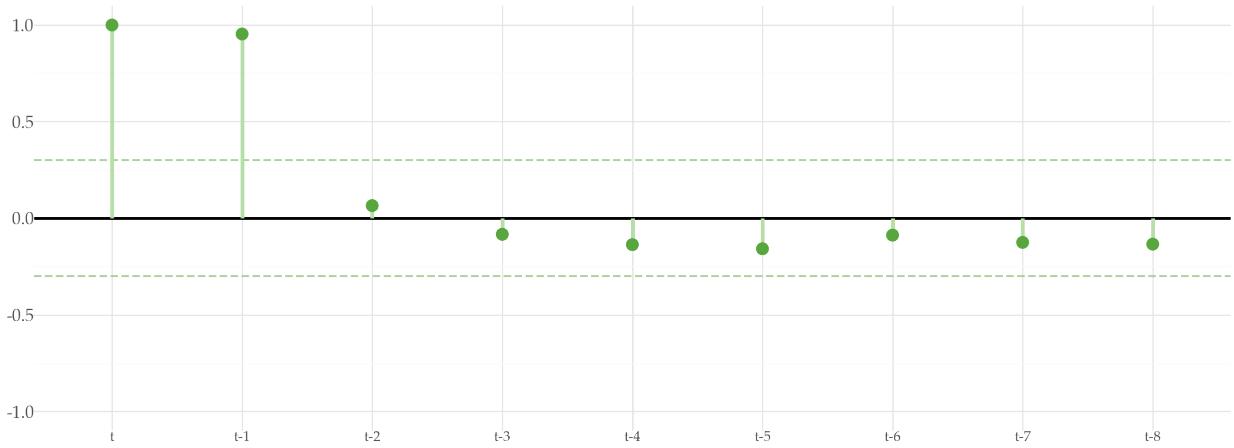

Figure 5: Partial Auto-correlation plot up to 8 lags. At each lag, the partial auto-correlation takes into account the previous correlations.

- The following lags show significant partial autocorrelation: t-1.
- None of the lags relative to the seasonal period (t-4 and t-8) show any significant partial autocorrelation.

Trend

Trend refers to the long-term change in the mean level of a time series. It reflects systematic and gradual changes in the data over time. Understanding the trend is important for identifying long-term growth or decline, structural changes, and making informed modeling decisions. This section examines the characteristics of the trend of the time series.

Trend Line Plot

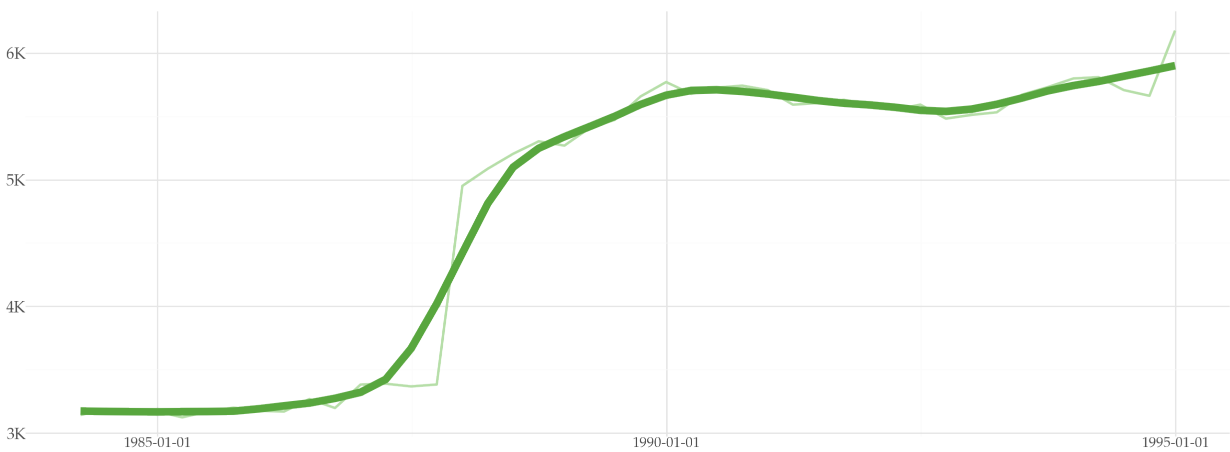

Figure 6: Time series trend plot.

- The time series has non-stationary trend according to the statistical test(s): KPSS, Augmented Dickey-Fuller, and Philips-Perron. There is a strong upward trend.
- The same tests were applied to analyse whether the time series is stationary around a constant level. The method(s) KPSS, Augmented Dickey-Fuller, and Philips-Perron reject this hypothesis.
- Including a trend explanatory variable which denotes the position of each observation does not improve forecasting performance.

Distribution of Differences

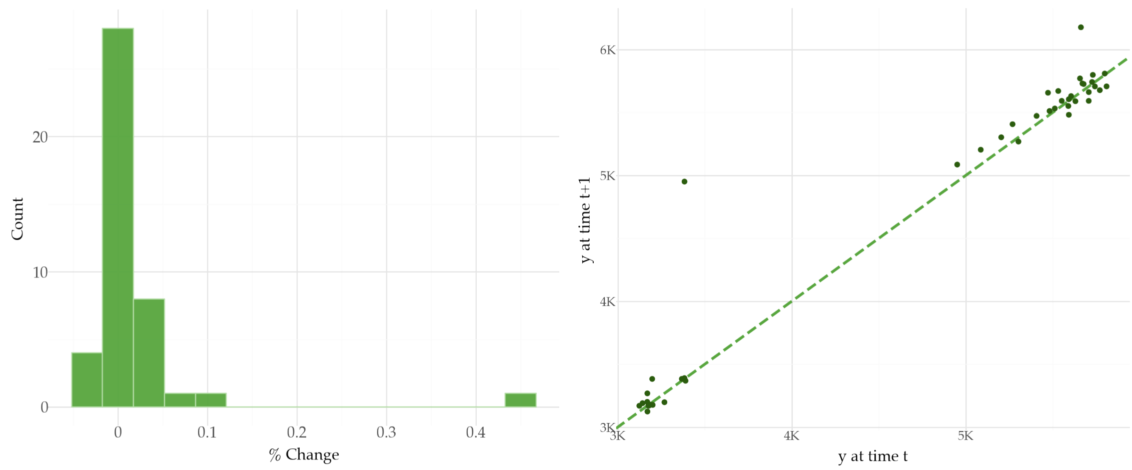

Figure 7: Distribution of percentage changes (left), and a Lag-plot (right). These plots help to understand how the data changes over consecutive observations. The histogram shown the distribution of these changes. The lag-plot depicts the randomness in the data. The time series shows greater randomness as the points deviate from the dotted line.

- The Kolmogorov-Smirnov test rejects the hypothesis that the differenced series is distributed according to the following distributions: Pareto, Chi-squared, Gaussian, and Power-law.
- The distribution with largest p-value is Cauchy (p-value equal to 0.62). But, we cannot reject the hypothesis that the differenced series follows the following distributions (ordered decreasingly by p-value): Log-Normal, Exponentially Modified Gaussian distribution, Logistic, Gamma, and Exponential.
- The excess kurtosis of the differenced series is equal to 25.51. This indicates that the data has a heavy tailed distribution.
- The skewness of the differenced series is equal to 4.86, indicates that the right tail is long relative to the left tail.
- **Forecasting experiments**: Taking first differences improves forecasting performance.

Seasonality

Seasonality represents recurring patterns or cycles that appear at regular intervals in time series data. These are predictable fluctuations that reflect periodic influences such as monthly, quarterly, or yearly cycles. Understanding seasonal patterns is crucial for forecasting, trend analysis, and identifying anomalies. This section examines the presence, strength, and characteristics of seasonal components in the input time series.

Seasonal Line Plot (Quarterly)

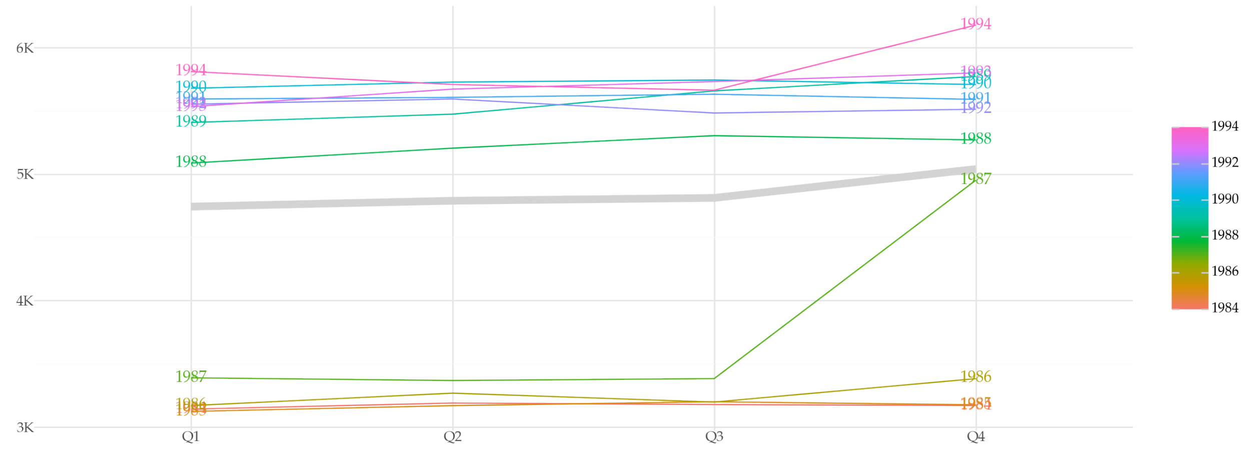

Figure 8: Seasonal plot of quarterly values grouped by year.

- All hypothesis tests carried out (Wang-Smith-Hyndman and OCSB) indicate that the time series is stationary in seasonality for a yearly period.
- **Forecasting experiments**: Including quarterly information in the predictive model decreases forecasting performance. This information was included as Fourier terms and repeating basis function terms in the explanatory variables.

Variance

Variance measures how data points spread around the average value in your time series. This section examines whether the variability remains stable (homoskedastic) or changes (heteroskedastic) over time. Understanding variance patterns is crucial for selecting appropriate modeling techniques, which can have a significant impact on forecasting accuracy.

Heteroskedasticity Testing

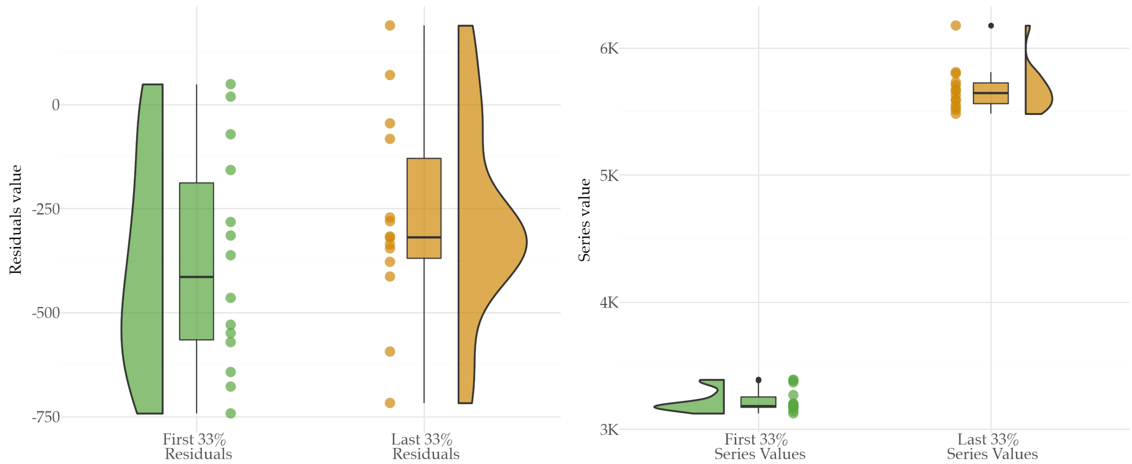

Figure 9: Time series residuals analysis. Difference in the distribution of the residuals (left) and the series (right) in the first and last thirds of the series, following a Goldfeld-Quand partition.

- In the analysis of seasonality we did not find significant differences in the dispersion among periodic groups of observations.
- The following tests suggest that the time series is heteroskedastic: Breusch-Pagan. But, other tests (White and Goldfeld-Quandt) fail to reject the hypothesis that the time series has a constant variance.
- **Forecasting experiments**: Transforming the series with the Box-Cox method improved forecasting performance. But, the transformation with the logarithm did not.
- No appropriate distribution was found to fit the data, before or after a Box-Cox transformation.

Change Detection

Change points denote significant shifts in the underlying distribution of time series. These structural changes can manifest as sudden shifts in level, trend, variance, or seasonal patterns. Detecting and understanding these points is crucial as they often indicate important events or regime changes that affect modeling decisions. This section identifies potential change points and assesses their impact on the overall analysis strategy.

Change Points

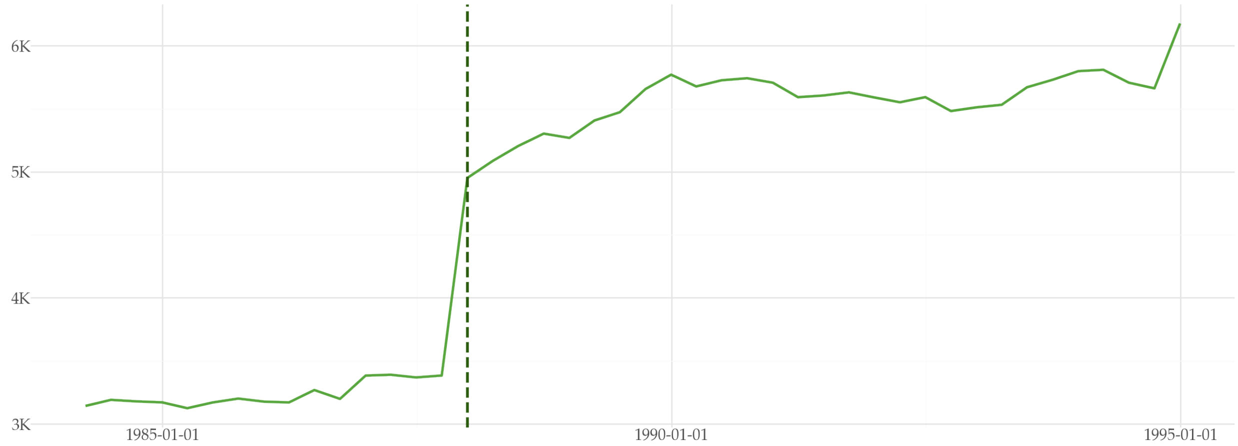

Figure 10: Time series plot with marked change points according to the PELT method.

- A single change point was found in the time series.
- The change point was found at 1987-52 where the time series shows an decreasing tendency.

Changes in Distribution

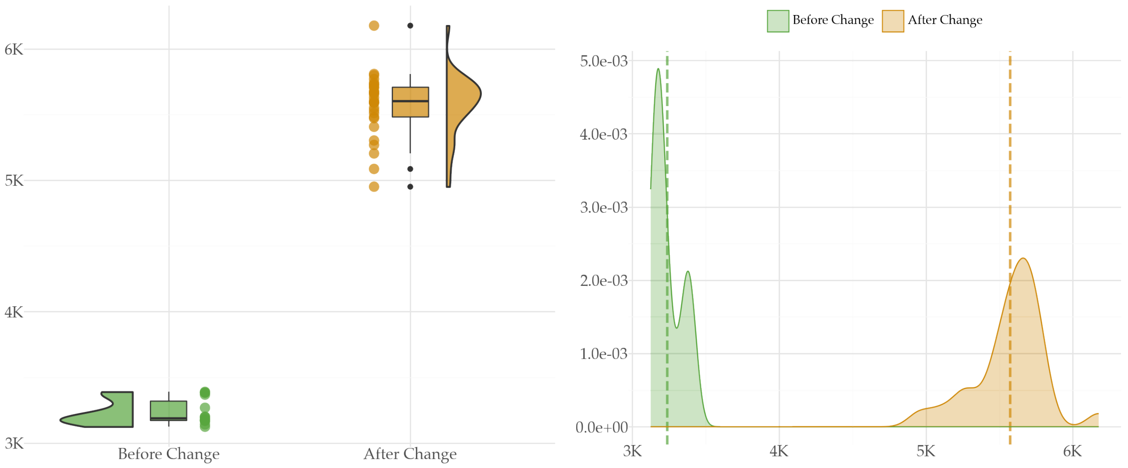

Figure 11: Time series analysis before and after the first detected change point occurs. Paired distributions before and after change (left), and overlapped density plot (right).

- The distribution before and after the first change point are significantly different.
- Before the change point, the data follows a Log-Normal distribution. But, after the first change point, a Logistic distribution is a better fit.